

hw5

1 Predicting ICU duration

Using the ICU cohort `mimiciv_icu_cohort.rds` you built in Homework 4, develop at least three machine learning approaches (logistic regression with enet regularization, random forest, boosting, SVM, MLP, etc) plus a model stacking approach for predicting whether a patient's ICU stay will be longer than 2 days. You should use the `los_long` variable as the outcome. Your algorithms can use patient demographic information (gender, age at ICU intime, marital status, race), ICU admission information (first care unit), the last lab measurements before the ICU stay, and first vital measurements during ICU stay as features. You are welcome to use any feature engineering techniques you think are appropriate; but make sure to not use features that are not available at an ICU stay's intime. For instance, `last_careunit` cannot be used in your algorithms.

1.1 library

```
sessionInfo()
```

R version 4.3.3 (2024-02-29)

Platform: x86_64-apple-darwin20 (64-bit)

Running under: macOS 15.3.1

Matrix products: default

BLAS: /Library/Frameworks/R.framework/Versions/4.3-x86_64/Resources/lib/libRblas.0.dylib

LAPACK: /Library/Frameworks/R.framework/Versions/4.3-x86_64/Resources/lib/libRlapack.dylib; LAPACK version 3.11.0

locale:

[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

time zone: America/Los_Angeles

tzcode source: internal

attached base packages:

[1] stats graphics grDevices utils datasets methods base

loaded via a namespace (and not attached):

[1] htmlwidgets_1.6.4 compiler_4.3.3 fastmap_1.2.0 cli_3.6.3
[5] tools_4.3.3 htmltools_0.5.8.1 rstudioapi_0.17.0 yaml_2.3.10
[9] rmarkdown_2.28 knitr_1.49 jsonlite_1.8.9 xfun_0.48
[13] digest_0.6.37 rlang_1.1.4 evaluate_1.0.3

```
library(tidyverse)
```

```

— Attaching core tidyverse packages — tidyverse 2.0.0 —
✓ dplyr      1.1.4      ✓ readr      2.1.5
✓ forcats    1.0.0      ✓ stringr    1.5.1
✓ ggplot2    3.5.1      ✓ tibble     3.2.1
✓ lubridate  1.9.3      ✓ tidyr      1.3.1
✓ purrr      1.0.2

— Conflicts — tidyverse_conflicts() —
✖ dplyr::filter() masks stats::filter()
✖ dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

```

```

library(ggplot2)
library(corrplot)

```

corrplot 0.95 loaded

```

library(lubridate)
library(miceRanger)
library(dplyr)
library(GGally)

```

Registered S3 method overwritten by 'GGally':
 method from
 +.gg ggplot2

```

library(gtsummary)
library(tidymodels)

```

```

— Attaching packages — tidymodels 1.2.0 —
✓ broom      1.0.7      ✓ rsample     1.2.1
✓ dials      1.4.0      ✓ tune        1.2.1
✓ infer      1.0.7      ✓ workflows   1.1.4
✓ modeldata  1.4.0      ✓ workflowsets 1.1.0
✓ parsnip    1.3.0      ✓ yardstick   1.3.2
✓ recipes    1.1.1

— Conflicts — tidymodels_conflicts() —
✖ scales::discard() masks purrr::discard()
✖ dplyr::filter()   masks stats::filter()
✖ recipes::fixed()  masks stringr::fixed()
✖ dplyr::lag()      masks stats::lag()
✖ yardstick::spec() masks readr::spec()
✖ recipes::step()   masks stats::step()
• Learn how to get started at https://www.tidymodels.org/start/

```

```

library(yardstick)
library(dials)
library(kernlab)

```

Attaching package: 'kernlab'

The following object is masked from 'package:dials':

buffer

The following object is masked from 'package:scales':

alpha

The following object is masked from 'package:purrr':

cross

The following object is masked from 'package:ggplot2':

alpha

```
library(ggthemes)
library(naniar)
library(kableExtra)
```

Attaching package: 'kableExtra'

The following object is masked from 'package:dplyr':

group_rows

```
library(stacks)
library(vip)
```

Attaching package: 'vip'

The following object is masked from 'package:utils':

vi

```
library(h2o)
```

Your next step is to start H2O:

```
> h2o.init()
```

For H2O package documentation, ask for help:

> ??h2o

After starting H2O, you can use the Web UI at <http://localhost:54321>
For more information visit <https://docs.h2o.ai>

Attaching package: 'h2o'

The following objects are masked from 'package:lubridate':

day, hour, month, week, year

The following objects are masked from 'package:stats':

cor, sd, var

The following objects are masked from 'package:base':

&&, %*%, %in%, ||, apply, as.factor, as.numeric, colnames,
colnames<-, ifelse, is.character, is.factor, is.numeric, log,
log10, log1p, log2, round, signif, trunc

```
h2o.init(nthreads = -1, max_mem_size = "8G")
```

Connection successful!

R is connected to the H2O cluster:

```
H2O cluster uptime:      6 hours 3 minutes
H2O cluster timezone:    America/Los_Angeles
H2O data parsing timezone: UTC
H2O cluster version:     3.44.0.3
H2O cluster version age:  1 year, 2 months and 27 days
H2O cluster name:        H2O_started_from_R_guojiayi_rnw010
H2O cluster total nodes: 1
H2O cluster total memory: 6.76 GB
H2O cluster total cores: 8
H2O cluster allowed cores: 8
H2O cluster healthy:     TRUE
H2O Connection ip:       localhost
H2O Connection port:     54321
H2O Connection proxy:    NA
H2O Internal Security:   FALSE
R Version:                R version 4.3.3 (2024-02-29)
```

Warning in h2o.clusterInfo():

Your H2O cluster version is (1 year, 2 months and 27 days) old. There may be a newer version available.

Please download and install the latest version from: https://h2o-release.s3.amazonaws.com/h2o/latest_stable.html

1.2 Data preprocessing and feature engineering.

```
icu_cohort <- readRDS("../hw4/mimiciv_shiny/mimiciv_icu_cohort.rds")
icu_cohort |>
  print(width = Inf)
```

A tibble: 94,458 × 44

	subject_id	hadm_id	stay_id.x	
	<dbl>	<dbl>	<dbl>	
1	10000032	29079034	39553978	
2	10000690	25860671	37081114	
3	10000980	26913865	39765666	
4	10001217	24597018	37067082	
5	10001217	27703517	34592300	
6	10001725	25563031	31205490	
7	10001843	26133978	39698942	
8	10001884	26184834	37510196	
9	10002013	23581541	39060235	
10	10002114	27793700	34672098	
	first_careunit			
	<fct>			
1	Medical Intensive Care Unit (MICU)			
2	Medical Intensive Care Unit (MICU)			
3	Medical Intensive Care Unit (MICU)			
4	Surgical Intensive Care Unit (SICU)			
5	Surgical Intensive Care Unit (SICU)			
6	Medical/Surgical Intensive Care Unit (MICU/SICU)			
7	Medical/Surgical Intensive Care Unit (MICU/SICU)			
8	Medical Intensive Care Unit (MICU)			
9	Cardiac Vascular Intensive Care Unit (CVICU)			
10	Coronary Care Unit (CCU)			
	last_careunit		intime	
	<fct>		<dtm>	
1	Medical Intensive Care Unit (MICU)		2180-07-23 14:00:00	
2	Medical Intensive Care Unit (MICU)		2150-11-02 19:37:00	
3	Medical Intensive Care Unit (MICU)		2189-06-27 08:42:00	
4	Surgical Intensive Care Unit (SICU)		2157-11-20 19:18:02	
5	Surgical Intensive Care Unit (SICU)		2157-12-19 15:42:24	
6	Medical/Surgical Intensive Care Unit (MICU/SICU)		2110-04-11 15:52:22	
7	Medical/Surgical Intensive Care Unit (MICU/SICU)		2134-12-05 18:50:03	
8	Medical Intensive Care Unit (MICU)		2131-01-11 04:20:05	
9	Cardiac Vascular Intensive Care Unit (CVICU)		2160-05-18 10:00:53	
10	Coronary Care Unit (CCU)		2162-02-17 23:30:00	
	outtime	los	admittime	disctime
	<dtm>	<dbl>	<dtm>	<dtm>
1	2180-07-23 23:50:47	0.410	2180-07-23 12:35:00	2180-07-25 17:55:00

2	2150-11-06	17:03:17	3.89	2150-11-02	18:02:00	2150-11-12	13:45:00
3	2189-06-27	20:38:27	0.498	2189-06-27	07:38:00	2189-07-03	03:00:00
4	2157-11-21	22:08:00	1.12	2157-11-18	22:56:00	2157-11-25	18:00:00
5	2157-12-20	14:27:41	0.948	2157-12-18	16:58:00	2157-12-24	14:55:00
6	2110-04-12	23:59:56	1.34	2110-04-11	15:08:00	2110-04-14	15:00:00
7	2134-12-06	14:38:26	0.825	2134-12-05	00:10:00	2134-12-06	12:54:00
8	2131-01-20	08:27:30	9.17	2131-01-07	20:39:00	2131-01-20	05:15:00
9	2160-05-19	17:33:33	1.31	2160-05-18	07:45:00	2160-05-23	13:30:00
10	2162-02-20	21:16:27	2.91	2162-02-17	22:32:00	2162-03-04	15:16:00
deathtime		admission_type		admit_provider_id			
<dtm>		<fct>		<chr>			
1	NA	EW EMER.		P060TX			
2	NA	EW EMER.		P26QQ4			
3	NA	EW EMER.		P060TX			
4	NA	EW EMER.		P3610N			
5	NA	Other		P2760U			
6	NA	EW EMER.		P32W56			
7	2134-12-06	12:54:00	URGENT	P67ATB			
8	2131-01-20	05:15:00	OBSERVATION ADMIT	P49AFC			
9	NA	SURGICAL SAME DAY ADMISSION		P8286C			
10	NA	OBSERVATION ADMIT		P46834			
admission_location		discharge_location		insurance	language	marital_status	
<fct>		<fct>		<chr>	<chr>	<chr>	
1	EMERGENCY ROOM	HOME		Medicaid	English	WIDOWED	
2	EMERGENCY ROOM	Other		Medicare	English	WIDOWED	
3	EMERGENCY ROOM	HOME HEALTH CARE		Medicare	English	MARRIED	
4	EMERGENCY ROOM	HOME HEALTH CARE		Private	Other	MARRIED	
5	PHYSICIAN REFERRAL	HOME HEALTH CARE		Private	Other	MARRIED	
6	Other	HOME		Private	English	MARRIED	
7	TRANSFER FROM HOSPITAL	DIED		Medicare	English	SINGLE	
8	EMERGENCY ROOM	DIED		Medicare	English	MARRIED	
9	PHYSICIAN REFERRAL	HOME HEALTH CARE		Medicare	English	SINGLE	
10	PHYSICIAN REFERRAL	HOME HEALTH CARE		Medicaid	English	<NA>	
edregtime		edouttime		hospital_expire_flag		gender	
<dtm>		<dtm>		<dbl>		<chr>	
1	2180-07-23	05:54:00	2180-07-23	14:00:00	0		F
2	2150-11-02	11:41:00	2150-11-02	19:37:00	0		F
3	2189-06-27	06:25:00	2189-06-27	08:42:00	0		F
4	2157-11-18	17:38:00	2157-11-19	01:24:00	0		F
5	NA	NA		0		F	
6	NA	NA		0		F	
7	NA	NA		1		M	
8	2131-01-07	13:36:00	2131-01-07	22:13:00	1		F
9	NA	NA		0		F	
10	2162-02-17	19:35:00	2162-02-17	23:30:00	0		M
anchor_age		anchor_year	anchor_year_group	dod	stay_id.y	Bicarbonate	
<dbl>		<dbl>	<chr>	<date>	<dbl>	<dbl>	
1	52	2180	2014 - 2016	2180-09-09	39553978	25	
2	86	2150	2008 - 2010	2152-01-30	37081114	26	
3	73	2186	2008 - 2010	2193-08-26	39765666	21	
4	55	2157	2011 - 2013	NA	34592300	30	

	5	55	2157	2011 – 2013	NA	34592300	30
6	46	2110	2011 – 2013	NA	31205490	NA	
7	73	2131	2017 – 2019	2134-12-06	39698942	28	
8	68	2122	2008 – 2010	2131-01-20	37510196	30	
9	53	2156	2008 – 2010	NA	39060235	24	
10	56	2162	2020 – 2022	2162-12-11	34672098	18	

	Chloride	Creatinine	Glucose	Potassium	Sodium	Hematocrit	wbc	stay_id	HR
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<dbl>
1	95	0.7	102	6.7	126	41.1	6.9	39553978	91
2	100	1	85	4.8	137	36.1	7.1	37081114	79
3	109	2.3	89	3.9	144	27.3	5.3	39765666	77
4	104	0.5	87	4.1	142	37.4	5.4	34592300	96
5	104	0.5	87	4.1	142	37.4	5.4	34592300	96
6	98	NA	NA	4.1	139	NA	NA	31205490	86
7	97	1.3	131	3.9	138	31.4	10.4	39698942	118
8	88	1.1	141	4.5	130	39.7	12.2	37510196	38
9	102	0.9	288	3.5	137	34.9	7.2	39060235	80
10	NA	3.1	95	6.5	125	34.3	16.8	34672098	111

	NBPs	NBPd	RR	BT	age_intime	race	los_long
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<lgl>
1	84	48	24	98.7	52	WHITE	FALSE
2	107	63	23	97.7	86	WHITE	TRUE
3	150	77	23	98	76	BLACK	FALSE
4	167	95	11	97.6	55	WHITE	FALSE
5	167	95	11	97.6	55	WHITE	FALSE
6	73	56	19	97.7	46	WHITE	FALSE
7	112	71	17	97.9	76	WHITE	FALSE
8	180	12	16	98.1	77	BLACK	TRUE
9	104	70	14	97.2	57	Other	FALSE
10	112	80	20	97.9	56	Other	TRUE

i 94,448 more rows

check the missing value

```
# check variables with more than 10000 missing values
icu_cohort %>%
  select_if(colSums(is.na(icu_cohort)) > 10000) %>%
  colnames()
```

```
[1] "deathtime" "edregtime" "edouttime" "dod"
```

```
# keep only variables with less than 10000 missing values
icu_cohort_discard <- icu_cohort %>%
  select_if(colSums(is.na(icu_cohort)) <= 10000)
```

delete data collected after ICU intime

```
icu_cohort_selected <- icu_cohort_discard |>
  select(-c("stay_id.y", "stay_id.x", "outtime", "disctime",
```

```
      "hospital_expire_flag", "last_careunit"))
print(icu_cohort_selected)
```

A tibble: 94,458 × 34

	subject_id	hadm_id	first_careunit	intime	los
	<dbl>	<dbl>	<fct>	<dtm>	<dbl>
1	10000032	29079034	Medical Intensive Care Unit (M...	2180-07-23 14:00:00	0.410
2	10000690	25860671	Medical Intensive Care Unit (M...	2150-11-02 19:37:00	3.89
3	10000980	26913865	Medical Intensive Care Unit (M...	2189-06-27 08:42:00	0.498
4	10001217	24597018	Surgical Intensive Care Unit (...)	2157-11-20 19:18:02	1.12
5	10001217	27703517	Surgical Intensive Care Unit (...)	2157-12-19 15:42:24	0.948
6	10001725	25563031	Medical/Surgical Intensive Car...	2110-04-11 15:52:22	1.34
7	10001843	26133978	Medical/Surgical Intensive Car...	2134-12-05 18:50:03	0.825
8	10001884	26184834	Medical Intensive Care Unit (M...	2131-01-11 04:20:05	9.17
9	10002013	23581541	Cardiac Vascular Intensive Car...	2160-05-18 10:00:53	1.31
10	10002114	27793700	Coronary Care Unit (CCU)	2162-02-17 23:30:00	2.91

i 94,448 more rows

i 29 more variables: admittime <dtm>, admission_type <fct>,
 # admit_provider_id <chr>, admission_location <fct>,
 # discharge_location <fct>, insurance <chr>, language <chr>,
 # marital_status <chr>, gender <chr>, anchor_age <dbl>, anchor_year <dbl>,
 # anchor_year_group <chr>, Bicarbonate <dbl>, Chloride <dbl>,
 # Creatinine <dbl>, Glucose <dbl>, Potassium <dbl>, Sodium <dbl>, ...

figure out missing value

```
# make a function to replace outliers to `NA`s
winsorize <- function(x, lower=0.01, upper=0.99) {
  qnt <- quantile(x, probs = c(lower, upper), na.rm = TRUE)
  x[x < qnt[1]] <- qnt[1]
  x[x > qnt[2]] <- qnt[2]
  x
}

# replace the extrem data with NA
icu_cohort_replace <- icu_cohort_selected %>%
  mutate(across(c("HR", "NBP", "NBPd", "RR", "BT",
                  "Bicarbonate", "Chloride", "Creatinine",
                  "Glucose", "Potassium", "Sodium",
                  "Hematocrit", "wbc"), winsorize))

# delete all rows with NA in los_long, marital_status
icu_cohort_replace <- icu_cohort_replace |>
  filter(!is.na(los_long)) |>
  filter(!is.na(marital_status)) |>
  # also delete labevent NA rows
  filter(!if_all(c("Bicarbonate", "Chloride", "Creatinine", "Glucose",
                  "Potassium", "Sodium", "Hematocrit", "wbc"), is.na))

# delete unnecessary columns
icu_cohort_replace <- icu_cohort_replace |>
  select(-c("admit_provider_id", "discharge_location")) |>
  print()
```



```
# A tibble: 83,043 × 32
  subject_id hadm_id first_careunit          intime          los
    <dbl>    <dbl> <fct>                <dtm>          <dbl>
1  10000032 29079034 Medical Intensive Care Unit (M... 2180-07-23 14:00:00 0.410
2  10000690 25860671 Medical Intensive Care Unit (M... 2150-11-02 19:37:00 3.89
3  10000980 26913865 Medical Intensive Care Unit (M... 2189-06-27 08:42:00 0.498
4  10001217 24597018 Surgical Intensive Care Unit (... 2157-11-20 19:18:02 1.12
5  10001217 27703517 Surgical Intensive Care Unit (... 2157-12-19 15:42:24 0.948
6  10001725 25563031 Medical/Surgical Intensive Car... 2110-04-11 15:52:22 1.34
7  10001843 26133978 Medical/Surgical Intensive Car... 2134-12-05 18:50:03 0.825
8  10001884 26184834 Medical Intensive Care Unit (M... 2131-01-11 04:20:05 9.17
9  10002013 23581541 Cardiac Vascular Intensive Car... 2160-05-18 10:00:53 1.31
10 10002155 20345487 Medical Intensive Care Unit (M... 2131-03-09 21:33:00 0.859
# i 83,033 more rows
# i 27 more variables: admittance <dtm>, admission_type <fct>,
# admission_location <fct>, insurance <chr>, language <chr>,
# marital_status <chr>, gender <chr>, anchor_age <dbl>, anchor_year <dbl>,
# anchor_year_group <chr>, Bicarbonate <dbl>, Chloride <dbl>,
# Creatinine <dbl>, Glucose <dbl>, Potassium <dbl>, Sodium <dbl>,
# Hematocrit <dbl>, wbc <dbl>, stay_id <int>, HR <dbl>, NBPs <dbl>, ...
```

1.2.1 keep cleaned data in a file

```
# fill in the NA value
if (file.exists("icu_cohort_filled.rds")){
  icu_cohort_filled <- read_rds("icu_cohort_filled.rds")
}else{
  imputed_data <- miceRanger(icu_cohort_replace,
                             m = 1,
                             max.depth = 8,
                             num.trees = 50)
  icu_cohort_filled <- completeData(imputed_data)
  icu_cohort_filled |>
    write_rds("icu_cohort_filled.rds")
}
icu_cohort_model <- as.data.frame(icu_cohort_filled) |>
  rename_with(~ gsub("^Dataset_1\\.\\.\"", "", .x)) |>
  # Keep necessary columns
  select(
    "los_long",
    "gender", "marital_status", "race", "first_careunit",
    "Bicarbonate", "Chloride", "Creatinine", "Glucose",
    "Potassium", "Sodium", "Hematocrit", "wbc",
    "HR", "NBPs", "NBPd", "RR", "BT",
    "age_intime", "subject_id", "stay_id", "hadm_id"
  ) |>
  mutate(los_long = factor(los_long, levels = c(FALSE, TRUE))) |>
  mutate(age_intime = as.numeric(age_intime)) |>
  mutate(gender = as.factor(gender)) |>
```

```
mutate(marital_status = as.factor(marital_status)) |>
mutate(race = as.factor(race))
```

print the summary table

```
summary_table <- icu_cohort_model %>%
  select("los_long",
    "gender", "marital_status", "race", "first_careunit",
    "Bicarbonate", "Chloride", "Creatinine", "Glucose",
    "Potassium", "Sodium", "Hematocrit", "wbc",
    "HR", "NBPs", "NBPd", "RR", "BT") |>
  tbl_summary(by = los_long)
summary_table
```

Characteristic	FALSE N = 43,044 ¹	TRUE N = 39,999 ¹
gender		
F	19,403 (45%)	17,488 (44%)
M	23,641 (55%)	22,511 (56%)
marital_status		
DIVORCED	3,413 (7.9%)	3,228 (8.1%)
MARRIED	20,465 (48%)	19,495 (49%)
SINGLE	13,621 (32%)	12,177 (30%)
WIDOWED	5,545 (13%)	5,099 (13%)
race		
ASIAN	1,470 (3.4%)	1,318 (3.3%)
BLACK	5,314 (12%)	4,747 (12%)
HISPANIC	1,857 (4.3%)	1,600 (4.0%)
Other	3,696 (8.6%)	3,858 (9.6%)
WHITE	30,707 (71%)	28,476 (71%)
first_careunit		
Cardiac Vascular Intensive Care Unit (CVICU)	6,747 (16%)	6,539 (16%)
Coronary Care Unit (CCU)	4,535 (11%)	4,589 (11%)

¹ n (%); Median (Q1, Q3)

Characteristic	FALSE N = 43,044 ¹	TRUE N = 39,999 ¹
Medical Intensive Care Unit (MICU)	9,909 (23%)	8,483 (21%)
Medical/Surgical Intensive Care Unit (MICU/SICU)	8,089 (19%)	6,002 (15%)
Neuro Intermediate	1,895 (4.4%)	3,137 (7.8%)
Surgical Intensive Care Unit (SICU)	5,913 (14%)	5,674 (14%)
Trauma SICU (TSICU)	4,670 (11%)	4,058 (10%)
Other	1,286 (3.0%)	1,517 (3.8%)
Bicarbonate	24.0 (21.0, 27.0)	24.0 (21.0, 27.0)
Chloride	102 (98, 105)	102 (98, 105)
Creatinine	1.00 (0.80, 1.40)	1.00 (0.80, 1.60)
Glucose	118 (98, 154)	121 (100, 158)
Potassium	4.20 (3.90, 4.60)	4.20 (3.90, 4.70)
Sodium	139.0 (136.0, 141.0)	138.0 (135.0, 141.0)
Hematocrit	36 (30, 40)	34 (29, 40)
wbc	9.0 (6.6, 12.6)	9.4 (6.8, 13.4)
HR	85 (74, 99)	87 (75, 102)
NBPs	122 (106, 139)	120 (104, 138)
NBPd	68 (58, 80)	67 (56, 79)
RR	18 (15, 22)	19 (15, 23)
BT	98.10 (97.60, 98.60)	98.20 (97.60, 98.70)
¹ n (%); Median (Q1, Q3)		

1.3 split the dataset

Partition data into 50% training set and 50% test set. Stratify partitioning according to `los_long`. For grading purpose, sort the data by `subject_id`, `hadm_id`, and `stay_id` and use the seed 203 for the initial data split. Below is the sample code.

```

set.seed(203)
# arrange the data
icu_cohort_model_split <- icu_cohort_model |>
  arrange(subject_id, hadm_id, stay_id) |>
  select(-c("subject_id",
            "hadm_id",
            "stay_id"))
data_split <- initial_split(
  icu_cohort_model_split,
  strata = "los_long",
  prop = 0.5
)
data_split

```

<Training/Testing/Total>
<41521/41522/83043>

```

icu_cohort_train <- training(data_split)
dim(icu_cohort_train)

```

[1] 41521 19

```

icu_cohort_test <- testing(data_split)
dim(icu_cohort_test)

```

[1] 41522 19

```
icu_cohort_train <- as.h2o(icu_cohort_train)
```

```

|
|
|
|=====| 100%

```

```
icu_cohort_test <- as.h2o(icu_cohort_test)
```

```

|
|
|
|=====| 100%

```

1.4 Logistic Regression

1.4.1 Model

```
logit_mod_h2o <- h2o.glm(
  x = c("gender", "marital_status", "race", "first_careunit",
        "Bicarbonate", "Chloride", "Creatinine", "Glucose",
        "Potassium", "Sodium", "Hematocrit", "wbc",
        "HR", "NBPs", "NBPd", "RR", "BT", "age_intime"),
  y = "los_long",
  training_frame = icu_cohort_train,
  family = "binomial",
  alpha = 0.5,
  lambda_search = TRUE,
  nfolds = 5,
  keep_cross_validation_predictions = TRUE
)
```

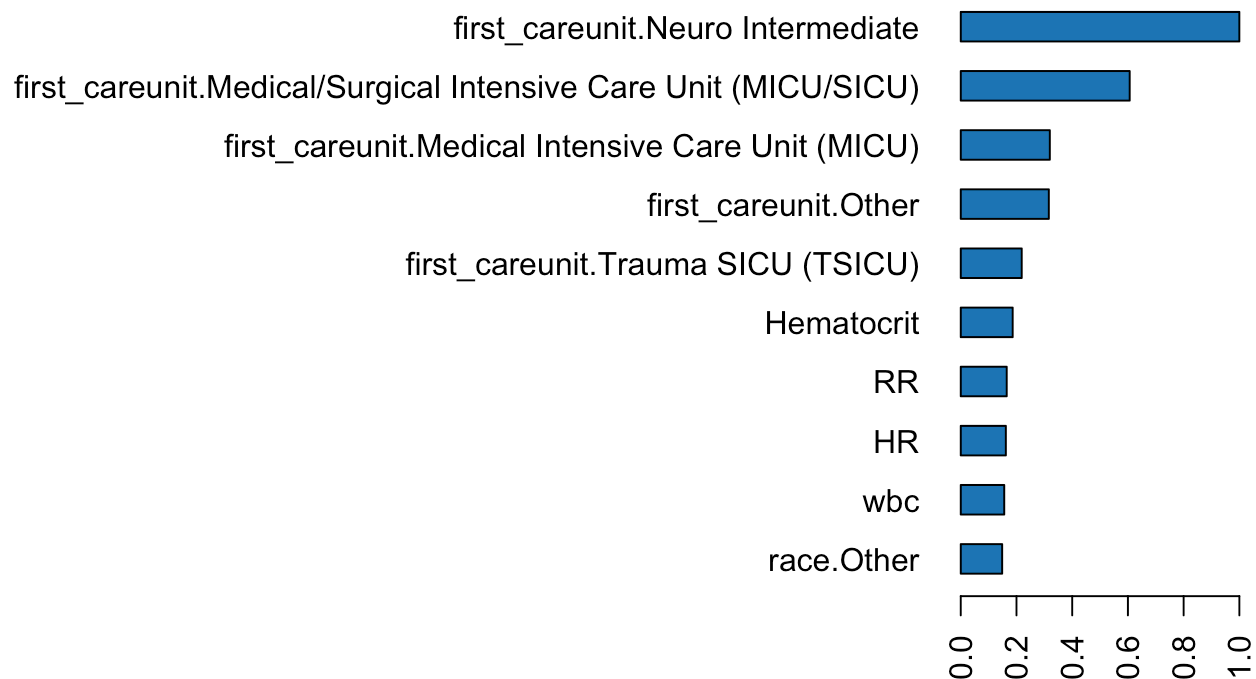
```
|
|
|
|=====| 20%
|
|=====| 38%
|
|=====| 100%
```

Warning in doTryCatch(return(expr), name, parentenv, handler): Reached maximum number of iterations 47!

1.4.2 Visualize CV results:

```
h2o.varimp_plot(logit_mod_h2o)
```

Variable Importance: GLM



```
h2o.performance(logit_mod_h2o)
```

H2O Binomial Metrics: glm

** Reported on training data. **

MSE: 0.2422779

RMSE: 0.4922174

LogLoss: 0.6774787

Mean Per-Class Error: 0.4933365

AUC: 0.5977081

AUCPR: 0.5695207

Gini: 0.1954162

R²: 0.02958268

Residual Deviance: 56259.18

AIC: 56311.18

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:

	FALSE	TRUE	Error	Rate
FALSE	473	21049	0.978022	=21049/21522
TRUE	173	19826	0.008650	=173/19999
Totals	646	40875	0.511115	=21222/41521

Maximum Metrics: Maximum metrics at their respective thresholds

	metric	threshold	value	idx
1	max f1	0.320247	0.651378	370
2	max f2	0.243312	0.822930	397
3	max f0point5	0.450966	0.560915	241
4	max accuracy	0.484125	0.571614	198
5	max precision	0.827308	1.000000	0
6	max recall	0.243312	1.000000	397
7	max specificity	0.827308	1.000000	0
8	max absolute_mcc	0.473170	0.143657	212
9	max min_per_class_accuracy	0.476334	0.570579	208
10	max mean_per_class_accuracy	0.473170	0.571850	212
11	max tns	0.827308	21522.000000	0
12	max fns	0.827308	19998.000000	0
13	max fps	0.225898	21522.000000	399
14	max tps	0.243312	19999.000000	397
15	max tnr	0.827308	1.000000	0
16	max fnr	0.827308	0.999950	0
17	max fpr	0.225898	1.000000	399
18	max tpr	0.243312	1.000000	397

Gains/Lift Table: Extract with ``h2o.gainsLift(<model>, <data>)`` or ``h2o.gainsLift(<model>, valid=<T/F>, xval=<T/F>)``

1.5 Random forest


1.5.1 Model

```
rf_mod_h2o <- h2o.randomForest(  
  x = c("gender", "marital_status", "race", "first_careunit",  
        "Bicarbonate", "Chloride", "Creatinine", "Glucose",  
        "Potassium", "Sodium", "Hematocrit", "wbc",  
        "HR", "NBPs", "NBPd", "RR", "BT", "age_intime"),  
  y = "los_long",  
  training_frame = icu_cohort_train,  
  ntrees = 100,  
  mtries = -1,  
  max_depth = 20,  
  min_rows = 5,  
  seed = 1234,  
  nfolds = 5,  
  balance_classes = TRUE,  
  keep_cross_validation_predictions = TRUE  
)
```

|
|
|

| 0%


=		1%
==		3%
====		6%
=====		11%
=====		15%
=====		18%
=====		20%
=====		23%
=====		28%
=====		32%
=====		35%
=====		37%
=====		40%
=====		45%
=====		50%
=====		52%
=====		54%
=====		57%
=====		62%
=====		65%
=====		67%
=====		70%
=====		72%
=====		76%
=====		80%
=====		83%



1.5.2 Visualize CV results

```
h2o.varimp_plot(rf_mod_h2o)
```

Variable Importance: DRF




1.6 XGBoost

1.6.1 Model

```
gb_mod_h2o <- h2o.gbm(
  x = c("gender", "marital_status", "race", "first_careunit",
```


```
    "Bicarbonate", "Chloride", "Creatinine", "Glucose",
    "Potassium", "Sodium", "Hematocrit", "wbc",
    "HR", "NBPs", "NBPd", "RR", "BT", "age_intime"),
y = "los_long",
training_frame = icu_cohort_train,
ntrees = 200,
max_depth = 6,
learn_rate = 0.05,
sample_rate = 0.8,
col_sample_rate = 0.8,
seed = 1234,
nfolds = 5,
keep_cross_validation_predictions = TRUE
)
```



|
=====| 100%

1.6.2 Visualize CV results

```
h2o.varimp_plot(gb_mod_h2o)
```



1.7 Model Stacking

```
feature_columns <- setdiff(names(icu_cohort_train), "los_long")

h2o_stacked_model <- h2o.stackedEnsemble(
  x = feature_columns,
  y = "los_long",
  training_frame = icu_cohort_train,
  base_models = list(rf_mod_h2o, gb_mod_h2o, logit_mod_h2o)
)
```

|
|
| 0%

```
|
|=====| 100%
```

1.7.1 Predict

```
h2o_stacked_predictions <- h2o.predict(h2o_stacked_model, icu_cohort_test) |>
  print()
```

```
|
|
|
|=====| 100%
predict      FALSE      TRUE
1      TRUE 0.5125255 0.4874745
2      TRUE 0.6031984 0.3968016
3     FALSE 0.6697046 0.3302954
4     FALSE 0.6697046 0.3302954
5      TRUE 0.5583365 0.4416635
6      TRUE 0.5873285 0.4126715
```

[41522 rows x 3 columns]

1.7.2 Model Performance

```
library(pROC)
```

Type 'citation("pROC")' for a citation.

Attaching package: 'pROC'

The following object is masked from 'package:h2o':

var

The following objects are masked from 'package:stats':

cov, smooth, var

```
h2o.performance(h2o_stacked_model, newdata = icu_cohort_test)
```

H2OBinomialMetrics: stackedensemble

MSE: 0.2352445

RMSE: 0.4850201

LogLoss: 0.6627929

Mean Per-Class Error: 0.463774

AUC: 0.6369982
 AUCPR: 0.6102181
 Gini: 0.2739963

Confusion Matrix (vertical: actual; across: predicted) for F1-optimal threshold:

	FALSE	TRUE	Error	Rate
FALSE	2573	18949	0.880448	=18949/21522
TRUE	942	19058	0.047100	=942/20000
Totals	3515	38007	0.479047	=19891/41522

Maximum Metrics: Maximum metrics at their respective thresholds

		metric threshold	value	idx
1		max f1	0.335106	0.657093 338
2		max f2	0.192474	0.822930 398
3		max f0point5	0.454596	0.585274 229
4		max accuracy	0.489978	0.598068 196
5		max precision	0.834882	1.000000 0
6		max recall	0.192474	1.000000 398
7		max specificity	0.834882	1.000000 0
8		max absolute_mcc	0.454596	0.196925 229
9	max min_per_class_accuracy	0.474861	0.596550	210
10	max mean_per_class_accuracy	0.454596	0.597682	229
11		max tns	0.834882	21522.000000 0
12		max fns	0.834882	19996.000000 0
13		max fps	0.187252	21522.000000 399
14		max tps	0.192474	20000.000000 398
15		max tnr	0.834882	1.000000 0
16		max fnr	0.834882	0.999800 0
17		max fpr	0.187252	1.000000 399
18		max tpr	0.192474	1.000000 398

Gains/Lift Table: Extract with `h2o.gainsLift(<model>, <data>)` or
 `h2o.gainsLift(<model>, valid=<T/F>, xval=<T/F>)`

```
# stacking model
perf_stack <- h2o.performance(h2o_stacked_model, newdata = icu_cohort_test)
auc_stack <- h2o.auc(perf_stack)

# rf model
perf_rf <- h2o.performance(rf_mod_h2o, newdata = icu_cohort_test)
auc_rf <- h2o.auc(perf_rf)

# gb model
perf_gb <- h2o.performance(gb_mod_h2o, newdata = icu_cohort_test)
auc_gb <- h2o.auc(perf_gb)

# logit model
perf_logit <- h2o.performance(logit_mod_h2o, newdata = icu_cohort_test)
auc_logit <- h2o.auc(perf_logit)
```

```
# table
auc_table <- data.frame(
  Model = c("Stacked Ensemble", "Random Forest",
            "GBM", "Logistic Regression"),
  AUC = c(auc_stack, auc_rf, auc_gb, auc_logit)
) |>
print()
```

	Model	AUC
1	Stacked Ensemble	0.6369982
2	Random Forest	0.6263728
3	GBM	0.6341979
4	Logistic Regression	0.6030615

conclusion The table above compares the AUC (Area Under the ROC Curve) values for four different classification models—Stacked Ensemble, Random Forest, GBM (Gradient Boosting Machine), and Logistic Regression. The Stacked Ensemble achieves the highest AUC at 0.6370, followed by GBM at 0.6342, Random Forest at 0.6264, and Logistic Regression at 0.6031. The stacked model slightly outperforms the individual algorithms.

1.7.3 Predicting compare

```
# Logistic Regression
perf_logit <- h2o.performance(logit_mod_h2o, newdata = icu_cohort_test)
cm_logit <- h2o.confusionMatrix(perf_logit, thresholds = 0.5)
```

Warning in h2o.find_row_by_threshold(object, t): Could not find exact threshold: 0.5 for this set of metrics; using closest threshold found: 0.500178129651075. Run `h2o.predict` and apply your desired threshold on a probability column.

```
cat("\n--- Logistic Regression Confusion Matrix (threshold=0.5) ---\n")
```

```
--- Logistic Regression Confusion Matrix (threshold=0.5) ---
```

```
print(cm_logit)
```

Confusion Matrix (vertical: actual; across: predicted) @ threshold = 0.500178129651075:

	FALSE	TRUE	Error	Rate
FALSE	14770	6752	0.313725	=6752/21522
TRUE	10891	9109	0.544550	=10891/20000
Totals	25661	15861	0.424907	=17643/41522

```
# Random Forest
perf_rf <- h2o.performance(rf_mod_h2o, newdata = icu_cohort_test)
cm_rf <- h2o.confusionMatrix(perf_rf, thresholds = 0.5)
```

Warning in h2o.find_row_by_threshold(object, t): Could not find exact threshold: 0.5 for this set of metrics; using closest threshold found: 0.499992401930058. Run `h2o.predict` and apply your desired threshold on a probability column.

```
cat("\n--- Random Forest Confusion Matrix (threshold=0.5) ---\n")
```

--- Random Forest Confusion Matrix (threshold=0.5) ---

```
print(cm_rf)
```

Confusion Matrix (vertical: actual; across: predicted) @ threshold = 0.499992401930058:

	FALSE	TRUE	Error	Rate
FALSE	14726	6796	0.315770	=6796/21522
TRUE	10198	9802	0.509900	=10198/20000
Totals	24924	16598	0.409277	=16994/41522

```
# GBM
```

```
perf_gb <- h2o.performance(gb_mod_h2o, newdata = icu_cohort_test)
cm_gb <- h2o.confusionMatrix(perf_gb, thresholds = 0.5)
```

Warning in h2o.find_row_by_threshold(object, t): Could not find exact threshold: 0.5 for this set of metrics; using closest threshold found: 0.500408734793412. Run `h2o.predict` and apply your desired threshold on a probability column.

```
cat("\n--- GBM Confusion Matrix (threshold=0.5) ---\n")
```

--- GBM Confusion Matrix (threshold=0.5) ---

```
print(cm_gb)
```

Confusion Matrix (vertical: actual; across: predicted) @ threshold = 0.500408734793412:

	FALSE	TRUE	Error	Rate
FALSE	14496	7026	0.326457	=7026/21522
TRUE	9744	10256	0.487200	=9744/20000
Totals	24240	17282	0.403882	=16770/41522

```
# Stacked Ensemble
```

```
perf_stack <- h2o.performance(h2o_stacked_model, newdata = icu_cohort_test)
cm_stack <- h2o.confusionMatrix(perf_stack, thresholds = 0.5)
```

Warning in h2o.find_row_by_threshold(object, t): Could not find exact threshold: 0.5 for this set of metrics; using closest threshold found:

0.499652532287864. Run `h2o.predict` and apply your desired threshold on a probability column.

```
cat("\n--- Stacked Ensemble Confusion Matrix (threshold=0.5) ---\n")
```

--- Stacked Ensemble Confusion Matrix (threshold=0.5) ---

```
print(cm_stack)
```

Confusion Matrix (vertical: actual; across: predicted) @ threshold = 0.499652532287864:

	FALSE	TRUE	Error	Rate
FALSE	14551	6971	0.323901	=6971/21522
TRUE	9751	10249	0.487550	=9751/20000
Totals	24302	17220	0.402726	=16722/41522

conclusion The Stacked Ensemble achieves the lowest overall error and the highest AUC, indicating that combining multiple base learners yields better predictive performance. Random Forest and GBM both outperform Logistic Regression, reflecting the benefit of more flexible, non-linear modeling. However, each model still shows substantial difficulty correctly identifying positive cases, as evidenced by higher error rates in the TRUE class. Overall, these results suggest that ensemble methods—especially stacking—can provide incremental yet meaningful improvements over individual algorithms.

1.7.4 ROC plot

```
true_labels <- as.vector(icu_cohort_test$los_long)
get_roc_df <- function(model, test_frame, true_labels, pos_class = "TRUE") {
  preds <- h2o.predict(model, test_frame)
  pred_prob <- as.vector(preds[[pos_class]])
  # use pROC
  roc_obj <- roc(true_labels, pred_prob)
  roc_df <- data.frame(
    fpr = 1 - rev(roc_obj$specificities),
    tpr = rev(roc_obj$sensitivities)
  )
  return(roc_df)
}


# get ROC for each model
roc_logit <- get_roc_df(logit_mod_h2o, icu_cohort_test, true_labels)
```

```
|
|
|
|=====| 100%
```

Setting levels: control = FALSE, case = TRUE

Setting direction: controls < cases

```
roc_rf <- get_roc_df(rf_mod_h2o, icu_cohort_test, true_labels)
```



Setting levels: control = FALSE, case = TRUE

Setting direction: controls < cases


```
roc_gb <- get_roc_df(gb_mod_h2o, icu_cohort_test, true_labels)
```



Setting levels: control = FALSE, case = TRUE

Setting direction: controls < cases

```
roc_stack <- get_roc_df(h2o_stacked_model, icu_cohort_test, true_labels)
```



Setting levels: control = FALSE, case = TRUE

Setting direction: controls < cases

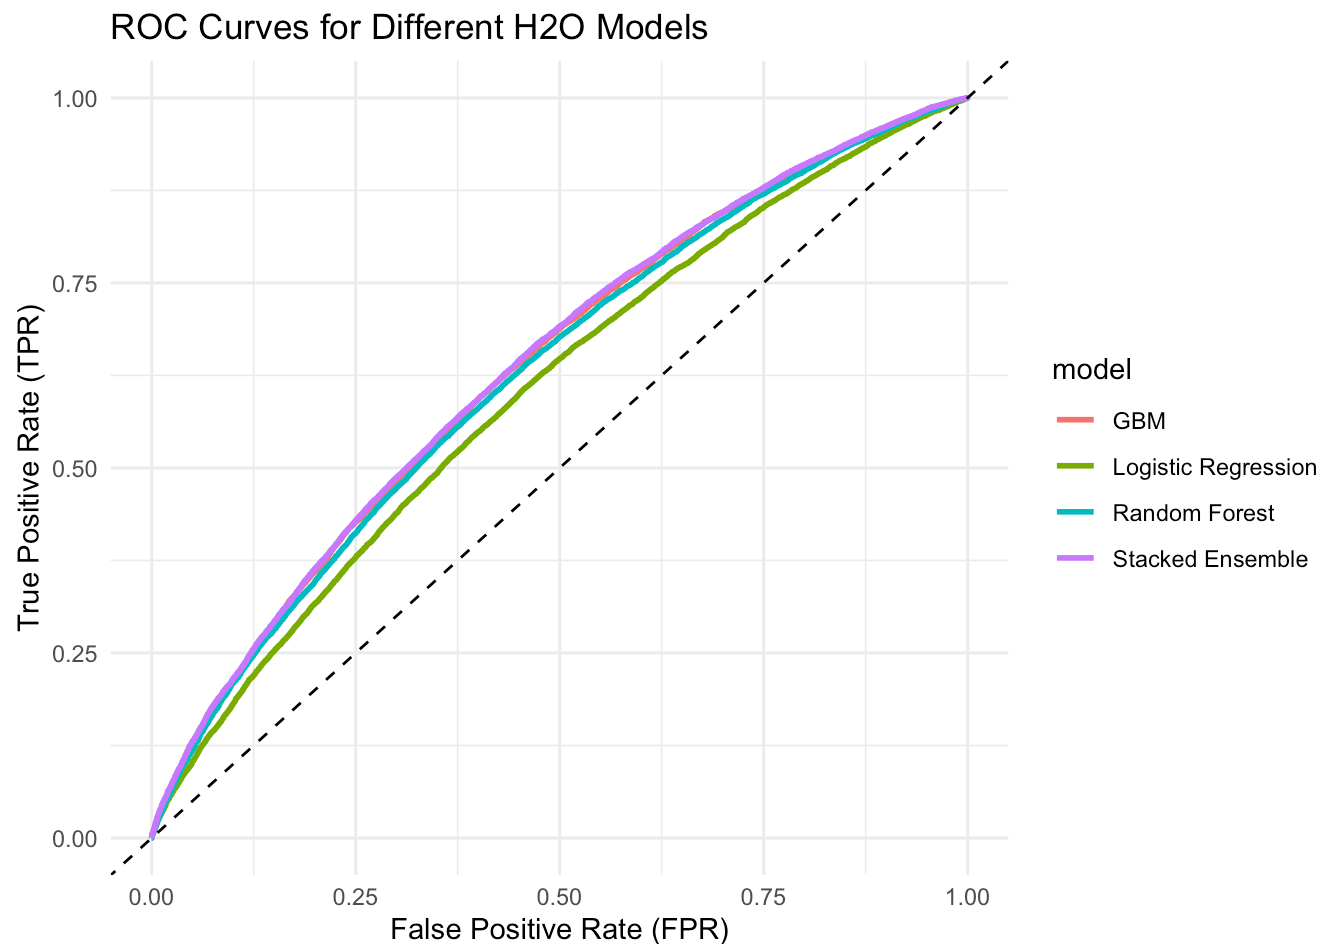
```
roc_logit$model <- "Logistic Regression"
roc_rf$model <- "Random Forest"
roc_gb$model <- "GBM"
roc_stack$model <- "Stacked Ensemble"

# merge all ROC data
roc_data <- rbind(roc_logit, roc_rf, roc_gb, roc_stack)

# ggplot
ggplot(roc_data, aes(x = fpr, y = tpr, color = model)) +
  geom_line(size = 1) +
  geom_abline(linetype = "dashed", color = "black") +
  labs(title = "ROC Curves for Different H2O Models",
       x = "False Positive Rate (FPR)",
```

```
y = "True Positive Rate (TPR)" +  
theme_minimal()
```

Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
i Please use `linewidth` instead.



conclusion From the figure it can be seen that stacking model has the best ROC curve performance, which is similar to that of GBM, random forests are slightly inferior to them, and logistic regression has the worst ROC performance.

What are the most important features in predicting long ICU stays? How do the models compare in terms of performance and interpretability? From the feature importance analysis, we can see all three models agree that first careunit is one of the most important features in predicting long ICU stays. The last lab measurements before the ICU stay and first vital measurements during ICU stay are also in top important features. This suggests that the last lab measurements before the ICU stay and first vital measurements during ICU stay are important indicators of the patient's condition and may be useful in predicting long ICU stays. In this case, the trade-off between performance and interpretability is clear. While the model stacking approach gives the best ROC AUC, it does so at the cost of interpretability. On the other hand, logistic regression offers ease of interpretation but doesn't perform as well. The gradient boosting model presents a good balance, with relatively high performance metrics and a degree of interpretability through feature importance scores.

